

helm部署OpenVPN 4.2.3

k8s部署openvpn打通k8s网络将Kubernetes集群网络暴露给本地开发网络

简单介绍:

VPN直译就是虚拟专用通道，是提供给企业之间或者个人与公司之间[安全数据传输](#)的隧道，OpenVPN无疑是Linux下开源VPN的先锋，提供了良好的性能和友好的用户GUI。

它大量使用了OpenSSL加密库中的SSLv3/TLSv1协议函数库。

目前OpenVPN能在Solaris、Linux、OpenBSD、FreeBSD、NetBSD、Mac OS X与Microsoft Windows以及Android和iOS上运行，并包含了许多安全性的功能。它并不是一个基于Web的VPN软件，也不与IPsec及其他VPN软件包兼容

官方仓库地址: [Kubeapps Hub](#)

客户端下载地址: [下载地址](#)

若该网址无法访问，可以到下面的链接进行下载

百度网盘: [点击直达](#) 提取码: l0qv

确认集群helm安装正确

```
[root@master01 ~]# helm version
version.BuildInfo{Version:"v3.8.2",
GitCommit:"6e3701edea09e5d55a8ca2aae03a68917630e91b", GitTreeState:"clean",
GoVersion:"go1.17.5"}
```

下载Chart包，修改values.yaml文件

```
helm repo add stable https://kubernetes-charts.storage.googleapis.com
helm pull stable/openvpn
```

或者

```
wget http://mirror.azure.cn/kubernetes/charts/openvpn-4.2.5.tgz
tar -zxvf openvpn-4.2.5.tgz
cd openvpn
```

修改配置文件,详情参考[stable/openvpn-github](#)

```
cp values.yaml openvpn-values.yaml
```

```
replicaCount: 1
updateStrategy: {}
imagePullSecretName:
image:
  repository: jfelten/openvpn-docker
```

```
tag: 1.1.0
pullPolicy: IfNotPresent
service:
  type: ClusterIP
  #externalPort: 443
  #internalPort: 443
  # hostPort: 443
  externalIPs: []
  nodePort: 32085
  # clusterIP: None
  # LoadBalancerSourceRanges: 0.0.0.0/0
  # loadBalancerIP: 10.0.0.1

  ## Here annotations can be added to the openvpn service
  # annotations:
  #   external-dns.alpha.kubernetes.io/hostname: vpn.example.com
  annotations: {}
  ## Here annotations can be added to the openvpn pod
  # podAnnotations:
  #   backup.ark.heptio.com/backup-volumes: certs
  podAnnotations: {}

  # Add privileged init container to enable IPv4 forwarding
  ipForwardInitContainer: true

resources:
  limits:
    cpu: 300m
    memory: 128Mi
  requests:
    cpu: 300m
    memory: 128Mi

readinessProbe:
  initialDelaySeconds: 5
  periodSeconds: 5
  successThreshold: 2

persistence:
  enabled: true

  storageClass: "nfs-client"
  accessMode: ReadWriteOnce
  size: 2M
openvpn:
  # Network allocated for openvpn clients (default: 10.240.0.0).
  OVPN_NETWORK: 10.240.0.0
  # Network subnet allocated for openvpn client (default: 255.255.0.0).
  OVPN_SUBNET: 255.255.0.0
  # Protocol used by openvpn tcp or udp (default: udp).
  OVPN_PROTO: udp
  OVPN_K8S_POD_NETWORK: "172.20.0.0" #k8s pod地址
  OVPN_K8S_POD_SUBNET: "255.255.0.0"
  OVPN_K8S_SVC_NETWORK: "172.21.0.0" #k8s svc地址
  OVPN_K8S_SVC_SUBNET: "255.255.0.0"

  DEFAULT_ROUTE_ENABLED: true
```

```
dhcptionDomain: true
# Redirect all client traffic through VPN
redirectGateway: true

useCr1: false
taKey: true
cipher: AES-256-CBC
istio:
  enabled: false
  proxy:
    port: 15001
iptablesExtra: []

ccd:
  enabled: false
  config: {}

nodeSelector: {}
tolerations: []
```

相关文档参考：

www.1nth.com/post/k8s-op...

zhuanlan.zhihu.com/p/491...

pythontaotao.github.io/2...

部署并进行验证

可以直接helm部署，也可以生成yaml文件部署

方法一：

```
kubectl create ns openvpn
helm install openvpn -n openvpn ./ -f openvpn-values.yaml
helm status openvpn -n openvpn
```

方法二：我才用此方法

```
helm install openvpn -n openvpn ./ -f openvpn-values.yaml --dry-run --debug
```

生成yaml，在进行部署

生成的文件

```
---
# Source: openvpn/templates/config-openvpn.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: openvpn
  labels:
    app: openvpn
    chart: openvpn-4.2.5
```

```

release: openvpn
heritage: Helm
data:
  setup-certs.sh: |-
    #!/bin/bash
    EASY_RSA_LOC="/etc/openvpn/certs"
    cd $EASY_RSA_LOC
    SERVER_CERT="${EASY_RSA_LOC}/pki/issued/server.crt"
    if [ -e "$SERVER_CERT" ]
    then
      echo "found existing certs - reusing"
      if [ ! -e ${EASY_RSA_LOC}/pki/ta.key ]
      then
        echo "generating missed ta.key"
        openvpn --genkey --secret ${EASY_RSA_LOC}/pki/ta.key
      fi
    else
      cp -R /usr/share/easy-rsa/* $EASY_RSA_LOC
      ./easyrsa init-pki
      echo "ca\n" | ./easyrsa build-ca nopass
      ./easyrsa build-server-full server nopass
      ./easyrsa gen-dh
      openvpn --genkey --secret ${EASY_RSA_LOC}/pki/ta.key
    fi

  newClientCert.sh: |-
    #!/bin/bash
    EASY_RSA_LOC="/etc/openvpn/certs"
    cd $EASY_RSA_LOC
    MY_IP_ADDR="$2"
    ./easyrsa build-client-full $1 nopass
    cat >${EASY_RSA_LOC}/pki/$1.ovpn <<EOF
    client
    nobind
    dev tun
    remote ${MY_IP_ADDR} 443 udp

    cipher AES-256-CBC

    redirect-gateway def1

    <key>
    `cat ${EASY_RSA_LOC}/pki/private/$1.key`
    </key>
    <cert>
    `cat ${EASY_RSA_LOC}/pki/issued/$1.crt`
    </cert>
    <ca>
    `cat ${EASY_RSA_LOC}/pki/ca.crt`
    </ca>
    <tls-auth>
    `cat ${EASY_RSA_LOC}/pki/ta.key`
    </tls-auth>
    key-direction 1
    EOF
    cat pki/$1.ovpn

```

```

revokeClientCert.sh: |-
    #!/bin/bash
    EASY_RSA_LOC="/etc/openvpn/certs"
    cd $EASY_RSA_LOC
    ./easysrsa revoke $1
    ./easysrsa gen-cr1
    cp ${EASY_RSA_LOC}/pki/cr1.pem ${EASY_RSA_LOC}
    chmod 644 ${EASY_RSA_LOC}/cr1.pem

configure.sh: |-
    #!/bin/sh

    cidr2mask() {
        # Number of args to shift, 255..255, first non-255 byte, zeroes
        set -- $(( 5 - ($1 / 8) )) 255 255 255 255 $(( (255 << (8 - ($1 % 8)))
& 255 )) 0 0 0
        [ $1 -gt 1 ] && shift "$1" || shift
        echo ${1-0}.${2-0}.${3-0}.${4-0}
    }

    cidr2net() {
        local i ip mask netOctets octets
        ip="${1%/*}"
        mask="${1#*/}"
        octets=$(echo "$ip" | tr '.' '\n')

        for octet in $octets; do
            i=$((i+1))
            if [ $i -le $(( mask / 8 )) ]; then
                netOctets="$netOctets.$octet"
            elif [ $i -eq $(( mask / 8 + 1 )) ]; then
                netOctets="$netOctets.$((((octet / ((256 / ((2**((mask %
8)))))))) * ((256 / ((2**((mask % 8))))))))"
            else
                netOctets="$netOctets.0"
            fi
        done

        echo ${netOctets#}
    }

/etc/openvpn/setup/setup-certs.sh

iptables -t nat -A POSTROUTING -s 10.240.0.0/255.255.0.0 -o eth0 -j
MASQUERADE
mkdir -p /dev/net
if [ ! -c /dev/net/tun ]; then
    mknod /dev/net/tun c 10 200
fi

if [ "$DEBUG" == "1" ]; then
    echo ===== ${OVPN_CONFIG} =====
    cat "${OVPN_CONFIG}"
    echo =====
fi

```

```

intAndIP="$(ip route get 8.8.8.8 | awk '/8.8.8.8/ {print $5 "-" $7}')"
int="${intAndIP%-*}"
ip="${intAndIP#*-}"
cidr="$(ip addr show dev "$int" | awk -vip="$ip" '($2 ~ ip) {print $2}')"

NETWORK="$(cidr2net $cidr)"
NETMASK="$(cidr2mask ${cidr#*/})"
DNS=$(cat /etc/resolv.conf | grep -v '^#' | grep nameserver | awk '{print $2}')
SEARCH=$(cat /etc/resolv.conf | grep -v '^#' | grep search | awk '{$1=""; print $0}')
FORMATTED_SEARCH=""
for DOMAIN in $SEARCH; do
    FORMATTED_SEARCH="${FORMATTED_SEARCH}push \"dhcp-option DOMAIN-SEARCH ${DOMAIN}\"\\n"
done
cp -f /etc/openvpn/setup/openvpn.conf /etc/openvpn/
sed 's|OVPN_K8S_SEARCH|'"${FORMATTED_SEARCH}"'|' -i /etc/openvpn/openvpn.conf
sed 's|OVPN_K8S_DNS|'"${DNS}"'|' -i /etc/openvpn/openvpn.conf
sed 's|NETWORK|'"${NETWORK}"'|' -i /etc/openvpn/openvpn.conf
sed 's|NETMASK|'"${NETMASK}"'|' -i /etc/openvpn/openvpn.conf

# exec openvpn process so it receives lifecycle signals
exec openvpn --config /etc/openvpn/openvpn.conf
openvpn.conf: |-
server 10.240.0.0 255.255.0.0
verb 3

key /etc/openvpn/certs/pki/private/server.key
ca /etc/openvpn/certs/pki/ca.crt
cert /etc/openvpn/certs/pki/issued/server.crt
dh /etc/openvpn/certs/pki/dh.pem

tls-auth /etc/openvpn/certs/pki/ta.key 0

cipher AES-256-CBC

key-direction 0
keepalive 10 60
persist-key
persist-tun

proto udp
port 443
dev tun0
status /tmp/openvpn-status.log

user nobody
group nogroup

push "route NETWORK NETMASK"

```

```
push "route 172.20.0.0 255.255.0.0"
```

```
push "route 172.21.0.0 255.255.0.0"
```

```
OVPN_K8S_SEARCH
```

```
push "dhcp-option DNS 172.21.16.10"
```

```
---
```

```
# Source: openvpn/templates/certs-pvc.yaml
```

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

```
metadata:
```

```
  name: openvpn
```

```
  labels:
```

```
    app: openvpn
```

```
    chart: openvpn-4.2.5
```

```
    release: openvpn
```

```
    heritage: Helm
```

```
spec:
```

```
  accessModes:
```

```
    - "ReadWriteOnce"
```

```
  resources:
```

```
    requests:
```

```
      storage: "2M"
```

```
  storageClassName: "nfs-client"
```

```
---
```

```
# Source: openvpn/templates/openvpn-service.yaml
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: openvpn
```

```
  labels:
```

```
    app: openvpn
```

```
    chart: openvpn-4.2.5
```

```
    release: openvpn
```

```
    heritage: Helm
```

```
spec:
```

```
  ports:
```

```
    - name: openvpn
```

```
      port: 443
```

```
      targetPort: 443
```

```
      protocol: UDP
```

```
  selector:
```

```
    app: openvpn
```

```
    release: openvpn
```

```
  type: ClusterIP
```

```
---
```

```
# Source: openvpn/templates/openvpn-deployment.yaml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: openvpn
```

```
  labels:
```

```

app: openvpn
chart: openvpn-4.2.5
release: openvpn
heritage: Helm
spec:
  replicas: 1
  selector:
    matchLabels:
      app: openvpn
      release: openvpn
  template:
    metadata:
      labels:
        app: openvpn
        release: openvpn
      annotations:
        checksum/config:
c6d7908b4a71d5de2ee5f823eac2522e8de7c15ab026fcc0f4149a270886df1f
    spec:
      initContainers:
        - args:
            - -C
            - sysctl -w net.ipv4.ip_forward=1
          command:
            - /bin/sh
          image: busybox:1.29
          imagePullPolicy: IfNotPresent
          name: sysctl
          resources:
            requests:
              cpu: 1m
              memory: 1Mi
          securityContext:
            privileged: true
      containers:
        - name: openvpn
          image: "jfelten/openvpn-docker:1.1.0"
          imagePullPolicy: IfNotPresent
          command: ["/etc/openvpn/setup/configure.sh"]
          ports:
            - containerPort: 443
              name: openvpn
          securityContext:
            capabilities:
              add:
                - NET_ADMIN
          readinessProbe:
            initialDelaySeconds: 5
            periodSeconds: 5
            successThreshold: 2
          exec:
            command:
              - nc
              - -u
              - -Z
              - 127.0.0.1
              - "443"
          resources:

```



```

requests:
  cpu: "300m"
  memory: "128Mi"
limits:
  cpu: "300m"
  memory: "128Mi"
volumeMounts:
- mountPath: /etc/openvpn/setup
  name: openvpn
  readOnly: false
- mountPath: /etc/openvpn/certs
  name: certs
  readOnly: false
volumes:
- name: openvpn
  configMap:
    name: openvpn
    defaultMode: 0775
- name: certs
  persistentVolumeClaim:
    claimName: openvpn

```

账号管理脚本

下面整理一下生成clientkey的脚本,这里面的信息就是上面部署完成后输出的信息,需要修改参数 KEY_NAME 和 SERVICE_IP,执行完成后会生成 wangwu.ovpn 的文件,客户端使用该文件作为连接凭证

create-client-key.sh 创建用户

```

$ cat create-client-key.sh

#!/bin/bash

POD_NAME=$(kubectl get pods --namespace "default" -l "app=openvpn,release=my-openvpn" -o jsonpath='{ .items[0].metadata.name }')
SERVICE_NAME=$(kubectl get svc --namespace "default" -l "app=openvpn,release=my-openvpn" -o jsonpath='{ .items[0].metadata.name }')
SERVICE_IP=10.169.68.142
KEY_NAME=wangwu
kubectl --namespace "default" exec -it "$POD_NAME"
/etc/openvpn/setup/newClientCert.sh "$KEY_NAME" "$SERVICE_IP"
kubectl --namespace "default" exec -it "$POD_NAME" cat
"/etc/openvpn/certs/pki/$KEY_NAME.ovpn" > "$KEY_NAME.ovpn"

$ chmod +x get_clientkey.sh && ./get_clientkey.sh

```

rm-client-key.sh 注销用户

```
#!/usr/bin/env bash

KEY_NAME=$1
POD_NAME=$(kubectl get pods -n "openvpn" -l "app=openvpn,release=openvpn" -o jsonpath='{.items[0].metadata.name}')
kubectl -n "openvpn" exec -it "$POD_NAME" /etc/openvpn/setup/revokeClientCert.sh $KEY_NAME
```

创建客户端密钥

```
[root@izwz98n1yl2yxjc2xpwt8dz openvpn]# kubectl --namespace "openvpn" get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
openvpn       LoadBalancer  172.19.77.8    119.23.xxx.xx  443:30924/TCP    8d

[root@izwz98n1yl2yxjc2xpwt8dz openvpn]# ./create-client-key.sh sunnyhaha
[root@izwz98n1yl2yxjc2xpwt8dz openvpn]# ll user-openvpn-20220402/sunnyhaha.ovpn
-rw-r--r-- 1 root root 7445 Apr  2 09:50 user-openvpn-20220402/sunnyhaha.ovpn
```

上面的命令会生成client.ovpn配置文件，因为前面用Helm部署openvpn时，创建的Service是ClusterIP类型的，生成的配置文件中的VPN服务地址会有错误，需要修改这个配置文件，将remote配置修改如下地址为k8s边缘节点的ip，端口为前面暴露的1194端口，同时删除redirect-gateway def1这行内容

解析k8s内部svc

部署完成之后有个坑,不能解析k8s内部svc

```
nslookup kubernetes.default.svc.cluster.local
服务器:  Unknown
Address:  192.168.1.1

DNS request timed out.
        timeout was 2 seconds.
DNS request timed out.
```

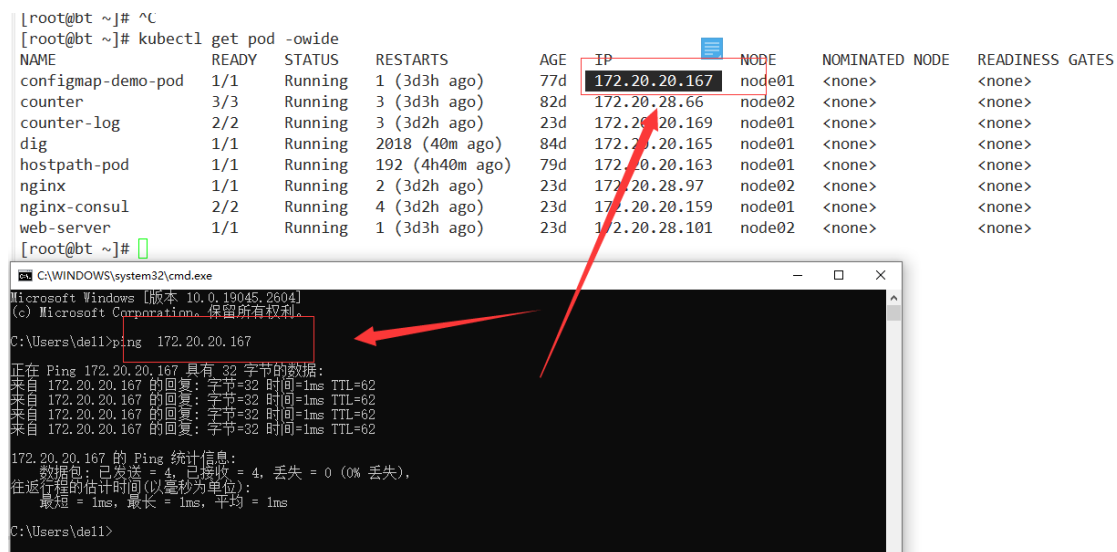
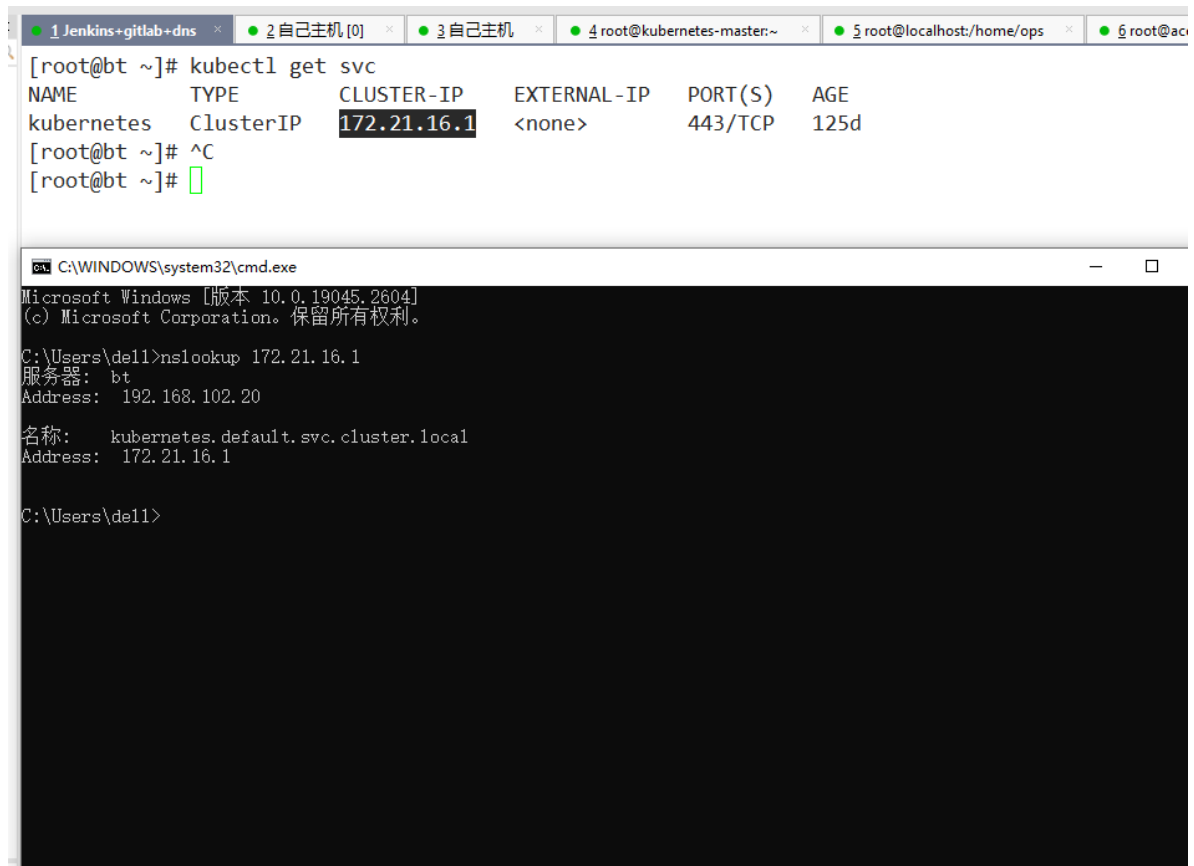
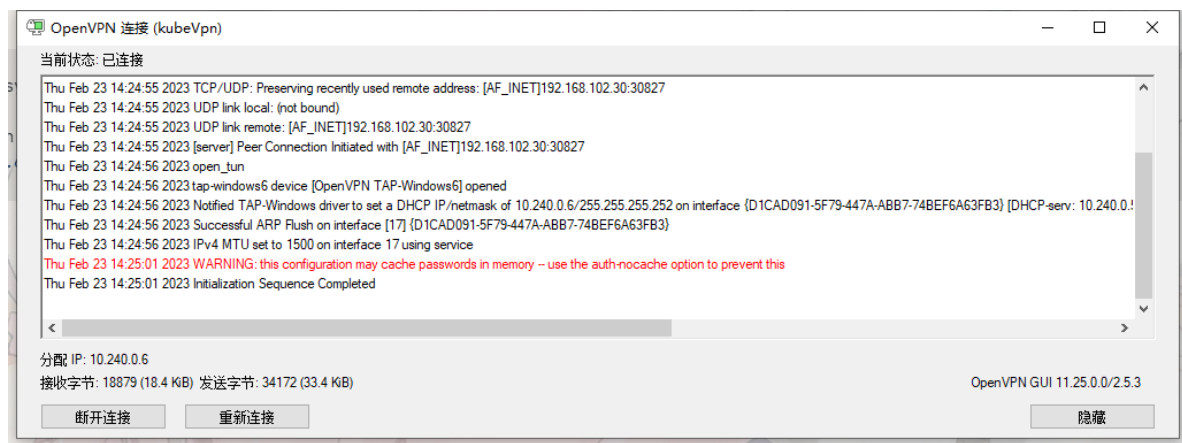
查看openvpn服务端配置

```
kubectl get cm -n openvpn openvpn -oyaml
.....
push "dhcp-option DNS OVPN_K8S_DNS"
```

因为我用的是node-local-dns,pod的DNS地址会变成PILLARLOCALDNS的默认地址169.254.20.10所以导致不能解析

修改服务端配置

```
kubectl get svc -n kube-system -l k8s-app=kube-dns -o jsonpath='{$.items[*].spec.clusterIP}'
10.68.0.2
kubectl get cm -n openvpn openvpn -oyaml
push "dhcp-option DNS 10.68.0.2"
```



```

sre          service-apollo-portal-server  <none>  apollo.ownit.top  192.168.102.40
[root@bt ~]# kubectl get ingress,svc,pod -n heian
NAME                                     CLASS    HOSTS                ADDRESS          PORTS    AGE
ingress.networking.k8s.io/ingress-heian  <none>   ingress.ownit.top    192.168.102.40  80      125d

NAME          TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/ingress-heian  ClusterIP   172.21.24.68   <none>         80/TCP     125d

NAME          READY   STATUS    RESTARTS      AGE
pod/ingress-heian-7ff9cbd9d5-2d2zq  1/1     Running   4 (3d2h ago)  86d
[root@bt ~]#
```

