

监控失败排查步骤

- 1、确认Service Monitor是否创建成功
- 2、确认Service Monitor标签是否配置正确
- 3、确认Prometheus是否生成了相关配置
- 4、确认存在Service Monitor匹配的Service
- 5、确认通过SVC能够访问程序的Metrics接口
- 6、确认SVC的端口跟Scheme和服务Monitor一致

一、以kube-controller-manager为例子

- 二进制部署的k8s+kube-Prometheus部署的监控默认没监控上
- 这也是监控kube-controller-manager的步骤了，同理kube-scheduler改一下端口即可

1、确认Service Monitor是否创建成功

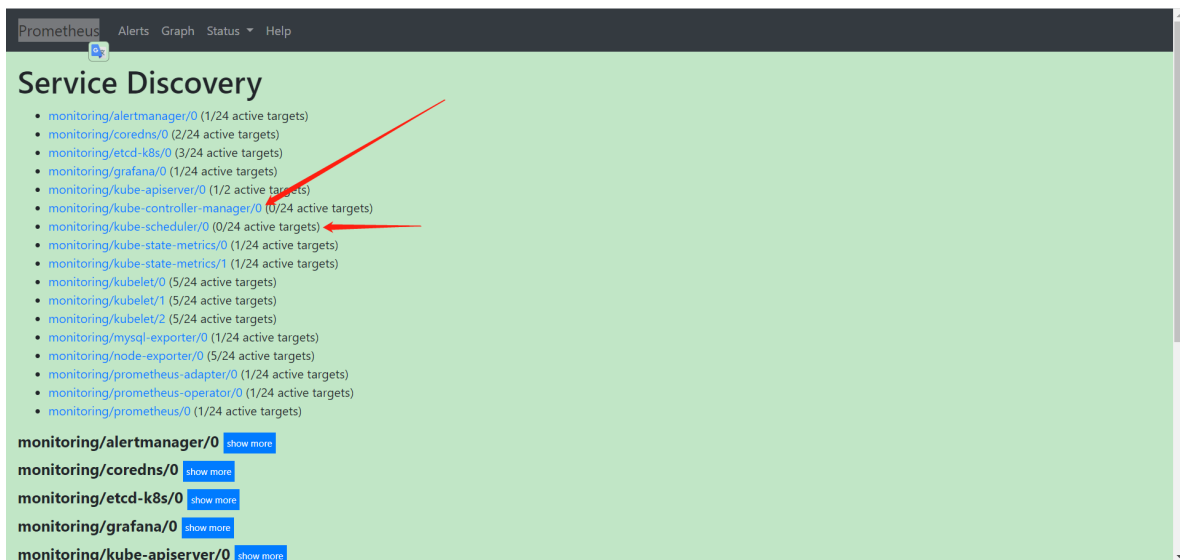
```
[root@k8s-master01 ~]# kubectl get servicemonitors -n monitoring | grep kube-  
controller-manager  
kube-controller-manager    6d4h
```

2、确认Service Monitor标签是否配置正确

```
[root@k8s-master01 ~]# kubectl get servicemonitors -n monitoring kube-  
controller-manager -oyaml | grep -A 3 ...  
...  
...  
  selector:  
    matchLabels:  
      k8s-app: kube-controller-manager  
...
```

3、确认Prometheus是否生成了相关配置

```
# 访问Prometheus的ui  
http://192.168.1.110:30062/service-discovery
```



4、确认存在Service Monitor匹配的Service

```
[root@k8s-master01 ~]# kubectl get servicemonitors -n monitoring kube-
controller-manager -oyaml | grep -A 3 ...
...
...
namespaceSelector:                                # 名称空间看namespaceSelector
  matchNames:
    - kube-system
selector:
  matchLabels:
    k8s-app: kube-controller-manager              # 匹配的是有这个标签svc,名称空间看
namespaceSelector
...

# 发现没有所以需要创建
[root@k8s-master01 ~]# kubectl get svc -n kube-system -l k8s-app=kube-
controller-manager
No resources found in kube-system namespace.

# 创建ep+svc
[root@k8s-master01 kube-controller-manager-监控]# cat kube-controller-
manager.yaml
apiVersion: v1
kind: Endpoints
metadata:
  labels:
    k8s-app: kube-controller-manager
    name: kube-controller-manage-monitor
    namespace: kube-system
subsets:
  - addresses:
    - ip: 192.168.1.110 # 改成master01宿主机的ip
    - ip: 192.168.1.111 # 改成master02宿主机的ip
    - ip: 192.168.1.112 # 改成master03宿主机的ip
    ports:
    - name: http-metrics
      port: 10252
      protocol: TCP
```

```

---
apiVersion: v1
kind: Service
metadata:
  labels:
    k8s-app: kube-controller-manager
  name: kube-controller-manage-monitor
  namespace: kube-system
spec:
  ports:
    - name: http-metrics
      port: 10252
      protocol: TCP
      targetPort: 10252
  sessionAffinity: None
  type: ClusterIP

```

5、确认通过kube-controller-manager SVC能够访问程序的Metrics接口

```

# 先确定kube-controller-manager服务是ok的
curl -k 127.0.0.1:10252/metrics

# 注意如果你的kube-controller-manager是监听在127.0.0.1上需要改成0.0.0.0
[root@k8s-master01 ~]# ss -ntlp | grep 10252
LISTEN  0  16384  127.0.0.1:10252  *:*  users:((("kube-
controller",pid=63861,fd=7))
# 更改方法【--address=0.0.0.0】
sed -i 's/address=127.0.0.1/address=0.0.0.0/' /usr/lib/systemd/system/kube-
controller-manager.service
systemctl daemon-reload && systemctl restart kube-controller-manager

# 查看svc
[root@k8s-master01 kube-controller-manager-监控]# kubectl get svc -n kube-system
kube-controller-manage-monitor

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
kube-controller-manage-monitor	ClusterIP	10.108.144.3	<none>

```

10252/TCP
# 再确认通过SVC能够访问程序的Metrics接口
curl -k 10.108.144.3:10252/metrics

```

6、确认kube-controller-manager SVC的端口跟Scheme和服务Monitor一致

```

[root@k8s-master01 ~]# kubectl get svc -n kube-system kube-controller-manage-
monitor -o yaml | grep -A 3 ports
--
ports:
- name: http-metrics
  port: 10252

```

protocol: TCP

此处不一样,所以要么跟etcd一样把证书挂载进去,要么修改kube-controller-manager的servicemonitors的port name,还有协议scheme一致

```
[root@k8s-master01 ~]# kubectl get servicemonitors -n monitoring kube-controller-manager -oyaml | grep port -A 1
    port: https-metrics
    scheme: https
```

直接edit

```
[root@k8s-master01 ~]# kubectl edit servicemonitors -n monitoring kube-controller-manager
[root@k8s-master01 ~]# kubectl get servicemonitors -n monitoring kube-controller-manager -oyaml | grep port -A 2
    port: http-metrics
    scheme: http
```

7、页面查看是否监控成功

- 看到以下数据说明监控是成功的

The image shows two screenshots of the Prometheus web interface. The top screenshot is the 'Service Discovery' page, which lists various monitored services and their target counts. A red arrow points to the 'monitoring/kube-controller-manager/0 (3/27 active targets)' entry. The bottom screenshot is the 'Alerts' page, showing a list of active alerts. A red arrow points to the 'Firing (4)' tab, with a red text label '这里告警也少了一个' (There is also one less alert here) next to it. The alerts list includes 'Watchdog (1 active)', 'CPUPhrottlingHigh (5 active)', and 'KubeSchedulerDown (1 active)'.

Service Discovery

- monitoring/alertmanager/0 (1/24 active targets)
- monitoring/coredns/0 (2/27 active targets)
- monitoring/etcd-k8s/0 (3/27 active targets)
- monitoring/grafana/0 (1/24 active targets)
- monitoring/kube-apiserver/0 (1/2 active targets)
- monitoring/kube-controller-manager/0 (3/27 active targets)
- monitoring/kube-scheduler/0 (0/27 active targets)
- monitoring/kube-state-metrics/0 (1/24 active targets)
- monitoring/kube-state-metrics/1 (1/24 active targets)
- monitoring/kubelet/0 (5/27 active targets)
- monitoring/kubelet/1 (5/27 active targets)
- monitoring/kubelet/2 (5/27 active targets)
- monitoring/mysql-exporter/0 (1/24 active targets)
- monitoring/node-exporter/0 (5/24 active targets)
- monitoring/prometheus-adapter/0 (1/24 active targets)
- monitoring/prometheus-operator/0 (1/24 active targets)
- monitoring/prometheus/0 (1/24 active targets)

Alerts

Inactive (99) Pending (0) **Firing (4)**

Watchdog (1 active)

CPUPhrottlingHigh (5 active)

KubeSchedulerDown (1 active)

