

# 产品告知业务不可用

# 1. 查看服务 发现node02的服务基本上属于不可用,但是却没有调度到node1 间接说明node02节点宿主机没问题 可能是网络出了问题

```
kubectl get pod -n ezhou-prod -o wide
NAME                                READY    STATUS              RESTARTS   AGE    IP              NODE                                NOMINATED NODE   READINESS GATES
-ddzcv                             1/1      Running             7          23h   173.0.3.196     k8s-node01.localdomain            <none>            <none>
-tcfvm                             1/1      Running             0          23h   10.36.1.204     k8s-node01.localdomain            <none>            <none>
595-tsjsxl                         2/2      Running             0          4d3h   173.0.4.128     k8s-node02.localdomain            <none>            <none>
6566d894b-vb7mf                   1/1      Running             10         23h   10.36.1.204     k8s-node01.localdomain            <none>            <none>
7dd5-ks7lc                         0/1      CrashLoopBackOff    53         21h   <none>          k8s-node02.localdomain            <none>            <none>
r                                  2/2      Running             2          39d   173.0.4.112     k8s-node02.localdomain            <none>            <none>
r                                  1/1      Running             8          23h   10.36.1.204     k8s-node01.localdomain            <none>            <none>
r                                  0/1      Running             8          34h   173.0.4.21      k8s-node02.localdomain            <none>            <none>
76f575fcd7-7kpfb                  0/1      Running             77         44d   173.0.4.40      k8s-node02.localdomain            <none>            <none>
5688cfcc84-dfrjt                  0/1      CrashLoopBackOff    87         44d   173.0.4.52      k8s-node02.localdomain            <none>            <none>
79c5457c-8zmkj                    0/1      Running             77         44d   173.0.4.23      k8s-node02.localdomain            <none>            <none>
696c8-rj4d7                        1/1      Running             80         44d   173.0.4.54      k8s-node02.localdomain            <none>            <none>
9m                                  1/1      Running             11         32h   173.0.3.211     k8s-node01.localdomain            <none>            <none>
8-flxcd                            0/1      RunContainerError    62         44d   <none>          k8s-node02.localdomain            <none>            <none>
8857-179kw                         1/1      Running             19         44d   173.0.4.111     k8s-node02.localdomain            <none>            <none>
df85-pm6c8                        1/1      Running             10         32h   173.0.3.210     k8s-node01.localdomain            <none>            <none>
-74cc56885d-khs24                 1/1      Running             0          23h   10.36.1.204     k8s-node01.localdomain            <none>            <none>
6f6-mz4g4                         0/1      Running             91         44d   173.0.4.17      k8s-node02.localdomain            <none>            <none>
66-zpqq1                          1/1      Running             10         22h   173.0.3.212     k8s-node01.localdomain            <none>            <none>
9-2jghv                           0/1      CrashLoopBackOff    69         44d   173.0.4.49      k8s-node02.localdomain            <none>            <none>
-6fcd455c69-8cd8s                 0/1      RunContainerError    92         44d   <none>          k8s-node02.localdomain            <none>            <none>
fd-crg2l                           1/1      Running             0          23h   10.36.1.204     k8s-node01.localdomain            <none>            <none>
7b-pxjb8                          1/1      Running             11         23h   10.36.1.204     k8s-node01.localdomain            <none>            <none>
78d5g-gh8j6                       2/2      Running             2          39d   173.0.4.106     k8s-node02.localdomain            <none>            <none>
566c9-5xs9j                       1/1      Running             22         44d   173.0.4.116     k8s-node02.localdomain            <none>            <none>
f9dd-zx46j                        0/1      RunContainerError    81         44d   <none>          k8s-node02.localdomain            <none>            <none>
z6                                  2/2      Running             0          34h   173.0.4.179     k8s-node02.localdomain            <none>            <none>
xl5k                               2/2      Running             4          45d   173.0.4.102     k8s-node02.localdomain            <none>            <none>
```

# 2. 登录node02查看,发现连接不上,但是可以telnet通跟ping通,尝试了几次ssh上了爆出如下问题,命令间接性执行失败 报bash:fork:retry

```
Last login: Fri Sep  2 15:28:25 2022 from 10.36.1.1
-bash: fork: retry: No child processes
[root@k8s-node02 ~]# ulimit -n
65535
```

# 3. 初步判断是进程数不够, 立即查看当前打开的

ps -A -ostat,ppid,pid,cmd | grep -e '^[Zz]'

```
[root@k8s-node02 ~]# ps h -Led -o user | sort | uniq -c|sort -rn
44844 root
233 1000
38 10001
36 472
7 polkitd
1 rpcuser
1 rpc
1 libstoragemgmt
1 dbus
1 chrony
```

# 4. 这进程数肯定不对啊,top看了下,发现有3w6的僵尸进程,所有我们ssh登录不上or连接上报错 bash:fork:retry,基本上可以判定是僵尸进程占用了进程数资源导致,新的应用or服务无法正常启动

```

[root@k8s-node02 ~]# top
top - 14:03:42 up 17 days, 19:25, 1 user, load average: 5.08, 2.71, 2.40
Tasks: 37077 total, 1 running, 628 sleeping, 0 stopped, 36197 zombie
%Cpu0  :  5.0 us,  5.5 sy,  0.0 ni, 89.4 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu1  :  2.7 us,  2.2 sy,  0.0 ni, 94.6 id,  0.0 wa,  0.0 hi,  0.4 si,  0.0 st
%Cpu2  :  2.7 us,  1.8 sy,  0.0 ni, 95.5 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu3  : 82.5 us,  0.0 sy,  0.0 ni, 17.5 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu4  :  4.0 us,  4.5 sy,  0.0 ni, 91.5 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu5  : 23.9 us,  4.9 sy,  0.0 ni, 71.2 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu6  :  5.9 us,  6.8 sy,  0.0 ni, 86.9 id,  0.0 wa,  0.0 hi,  0.5 si,  0.0 st
%Cpu7  : 81.1 us,  0.0 sy,  0.0 ni, 18.9 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu8  :  3.1 us,  1.3 sy,  0.0 ni, 95.5 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu9  :  1.4 us,  0.5 sy,  0.0 ni, 98.2 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu10 :  1.3 us,  1.3 sy,  0.0 ni, 97.3 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu11 :  2.7 us,  0.9 sy,  0.0 ni, 96.4 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu12 : 40.9 us,  5.8 sy,  0.0 ni, 53.3 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu13 :  6.2 us,  1.3 sy,  0.0 ni, 92.4 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu14 :  5.4 us,  1.3 sy,  0.0 ni, 93.3 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu15 : 12.2 us,  0.9 sy,  0.0 ni, 86.9 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu16 : 19.8 us,  6.8 sy,  0.0 ni, 73.0 id,  0.0 wa,  0.0 hi,  0.5 si,  0.0 st
%Cpu17 :  6.3 us,  5.9 sy,  0.0 ni, 87.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu18 : 24.9 us,  0.9 sy,  0.0 ni, 74.2 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu19 : 12.8 us,  0.9 sy,  0.0 ni, 85.8 id,  0.0 wa,  0.0 hi,  0.5 si,  0.0 st
%Cpu20 :  8.4 us,  0.4 sy,  0.0 ni, 91.2 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu21 :  6.2 us,  7.6 sy,  0.0 ni, 86.2 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu22 : 13.1 us,  6.6 sy,  0.0 ni, 79.9 id,  0.0 wa,  0.0 hi,  0.4 si,  0.0 st
%Cpu23 : 12.3 us,  0.9 sy,  0.0 ni, 86.3 id,  0.0 wa,  0.0 hi,  0.4 si,  0.0 st
%Cpu24 :  4.6 us,  3.2 sy,  0.0 ni, 92.2 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu25 : 13.9 us,  1.3 sy,  0.0 ni, 84.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu26 :  4.5 us,  3.2 sy,  0.0 ni, 73.0 id, 19.4 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu27 : 52.2 us,  3.0 sy,  0.0 ni, 44.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu28 :  2.2 us,  2.7 sy,  0.0 ni, 95.1 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu29 :  5.9 us,  1.4 sy,  0.0 ni, 92.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu30 :  0.4 us,  0.4 sy,  0.0 ni, 99.1 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
%Cpu31 : 23.1 us,  0.4 sy,  0.0 ni, 76.4 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st

```

# 过滤出僵尸进程

```
ps -A -ostat,ppid,pid,cmd | grep -e '^[zz]'
```

```
ps aux|grep pid # 查看父进程(主服务)
```

# 清理的话我是kill -9 pid

# 大量僵尸进程的的原因是甲方docker启动的服务没有做回收处理 可以通过bash/init/tini从而完成僵尸进程回收